

What is the N + 1 problem?

In simple terms, the N + 1 query problem happens when you:

1. Run 1 query to fetch a list of records
2. Then, for each of those N records, you run another query.

That's why it's called N + 1: one main query plus N additional queries.

Why It's a Problem

- Performance bottlenecks

Instead of running a single efficient query, you are now hitting the database multiple times. This slows your application significantly.

- Scalability issues

The more records you have in the database, severe it gets. With 10 users, you will run 11 queries. With 10,000 users you will run 10,001 queries — making the application unscalable.

Example of N + 1 Problem

Example of fetching all authors and their posts.

[github code solution](#)

Inefficient approach:

- Query 1: Get all the authors
- Query 2:
 - ... } For each author, fetch their posts
- Query N:

Example in Go with Gin

```
package controller

import (
    "net/http"
    // "database/sql"
    "github.com/gin-gonic/gin"

    "n-plus-one-problem/db"
    "n-plus-one-problem/models"
)

func GetAuthorsAndPosts(ctx *gin.Context) {
    // fetching all the authors
    var authors []models.Author
    query := "SELECT * FROM authors"
    if err := db.DB.Select(&authors, query); err != nil {
        ctx.JSON(http.StatusInternalServerError, gin.H{
            "error": "Error fetching the authors",
        })
    }
}
```

```

    })
    return
}

for i, author := range authors {
    var posts []models.Post
    query = "SELECT * FROM posts WHERE author_id = $1"
    if err := db.DB.Select(&posts, query, author.ID); err ==
nil {
        authors[i].Posts = posts
    }
}

ctx.JSON(http.StatusOK, authors)
}

```

Efficient approach:- Fetch all authors and their posts in one query using a SQL JOIN

```

package controller

import (
    "net/http"
    // "database/sql"
    "github.com/gin-gonic/gin"

    "n-plus-one-problem/db"
    "n-plus-one-problem/models"
)

func GetAuthorsAndPosts(ctx *gin.Context) {
    type AuthorWithPosts struct {
        AuthorID int    `db:"author_id" json:"author_id"`
        Name      string `db:"name" json:"name"`
        PostID    int    `db:"post_id" json:"post_id"`
    }
}

```

```

        Title      string `db:"title" json:"title"`
    }

    var results []AuthorWithPosts
    query := `
        SELECT a.id as author_id, a.name, p.id as post_id,
        p.title
        FROM authors a
        LEFT JOIN posts p ON p.author_id = a.id
    `

    if err := db.DB.Select(&results, query); err != nil {
        ctx.JSON(http.StatusInternalServerError, gin.H{
            "error": "Failed to fetch data",
        })
        return
    }

    ctx.JSON(http.StatusOK, results)
}

```

How to Avoid N + 1 Queries

1. Use SQL joins (JOIN, LEFT JOIN) instead of looping queries.
2. Leverage ORM features like:
 - Django: `.select_related`, `.prefetch_related`
 - SQLAlchemy: `joinedload`, `subqueryload`
 - Hibernate: `fetch joins`
3. Batch queries instead of per-record queries.

My conclusion

The N + 1 problem is sneaky — your code still works, but performance suffers silently. The solution is almost always about fetching smarter, not harder: reduce round trips, load related data in bulk, and let your database do the heavy lifting.